

■ Financial Document Q&A; Assistant - README

■ Project Overview

This project is a Streamlit-based web application that processes financial documents (PDF/Excel) and allows users to ask natural language questions about financial data. It uses Ollama (local LLM) to answer questions based on extracted document content.

■ Requirements

1. Python 3.9 or higher
2. Streamlit
3. Pandas
4. pdfplumber
5. Ollama (installed and running locally)
6. Llama2 model (pulled into Ollama using: `ollama pull llama2`)

■ Setup Instructions

1. Clone the repository or copy project files into a folder
2. Navigate into the folder:
`cd financial-qna-streamlit`
3. Create and activate virtual environment:
`python -m venv venv`
`venv\Scripts\activate` # On Windows
`source venv/bin/activate` # On Mac/Linux
4. Install dependencies:
`pip install -r requirements.txt`
5. Ensure Ollama is installed from <https://ollama.ai/download>
6. Pull the Llama2 model into Ollama:
`ollama pull llama2`

■ How to Run the Project (Commands)

Step 1: Start the Ollama server:
`ollama serve`

Step 2: Run the Streamlit app in a new terminal:
`streamlit run app.py`

■ Keep Ollama running in one terminal, and Streamlit in another.

■ Usage Instructions

1. Open the Streamlit app (URL shown in terminal, usually `http://localhost:8501`)
2. Upload a financial document (PDF or Excel)
3. Preview the extracted data in the interface

4. Ask questions like:
 - What is the total revenue?
 - What are the expenses for 2023?
 - Show me net profit
5. The app sends your question + extracted data to Ollama (Llama2) and displays an answer

■ Notes

- If you see 'No data extracted', check if the document has proper tables/text.
- Restart Ollama if you change models.
- Ollama and Streamlit must run together for the app to work.

■ Success Criteria

- Able to upload PDF/Excel
- Extract financial data
- Ask questions and get answers from Llama2
- Clean, simple Streamlit interface